





Swiss Institute of  
Bioinformatics



Using

# Large Language Models for Biodata Exploration: From Theory to Practice



Ana Claudia Sima, Vincent Emonet,  
Tarcisio Mendes, Panayiotis Smeros

kru@sib.swis



# Schedule

- The theory : LLM basics
- Coffee break @ 10:30
- From Theory (closer to) Practice : LLM question answering over biodata •

Lunch @ 12:00

- Hands-on : RAG-based LLM application for biodata exploration

- Coffee break @ 15:00



## Who we are – Knowledge Representation Unit @ SIB

• Tarcisio Mendes, Ana Claudia Sima – KRU co-leads (Lausanne / Zurich) • Vincent

Emonet – Research Software Engineer, lead developer of [ExpasyGPT](#)

• Panayiotis Smeros, Research Scientist, broad experience in ML / LLMs for information retrieval, classification



# What to expect from today

- An overview of LLMs, basics on how they are trained and potential uses
- A deep dive into using LLMs for question answering with context -
  - From unstructured and structured sources (inputs)
  - Towards unstructured and structured answers (outputs)
- RAG-based application for interacting with biodata
  - Build your own mini “ExpasyGPT” !




## Part I. The Theory

What is a Large Language Model?

A (deep learning) model trained to predict the next word in a sentence

How is it trained?

- Self-supervised learning on HUGE amounts of text
- Learning to predict 1 word at a time



Source: <https://amitnness.com>

Large??



“Large” = number of parameters


- 8B, 70B, 200B...1 Trillion? (GPT<sub>4o</sub>)

“Large” = amount of text seen

- GPT<sub>4</sub>: trained on 45TB (13 Trillion tokens!) of data (text, code)

Source: [GKD: Generalized Knowledge Distillation for Auto-regressive Sequence Models](#)<sub>9</sub>

## Closed-source LLMs: The GPT family



**GPT- 4o** (“research preview”)  
Multimodal

**GPT- o1**  
Reasoning

**GPT- 4.1**

**GPT- 4.5**



- Originally based on GPT3.5
- A “chatty” version of LLM



Note: ChatGPT is *not* an LLM per-se, but a service with many “bells and whistles” leveraging an LLM

# The “art” of prompting

- Prompting = a guiding question / instruction given to the model to shape the generated response



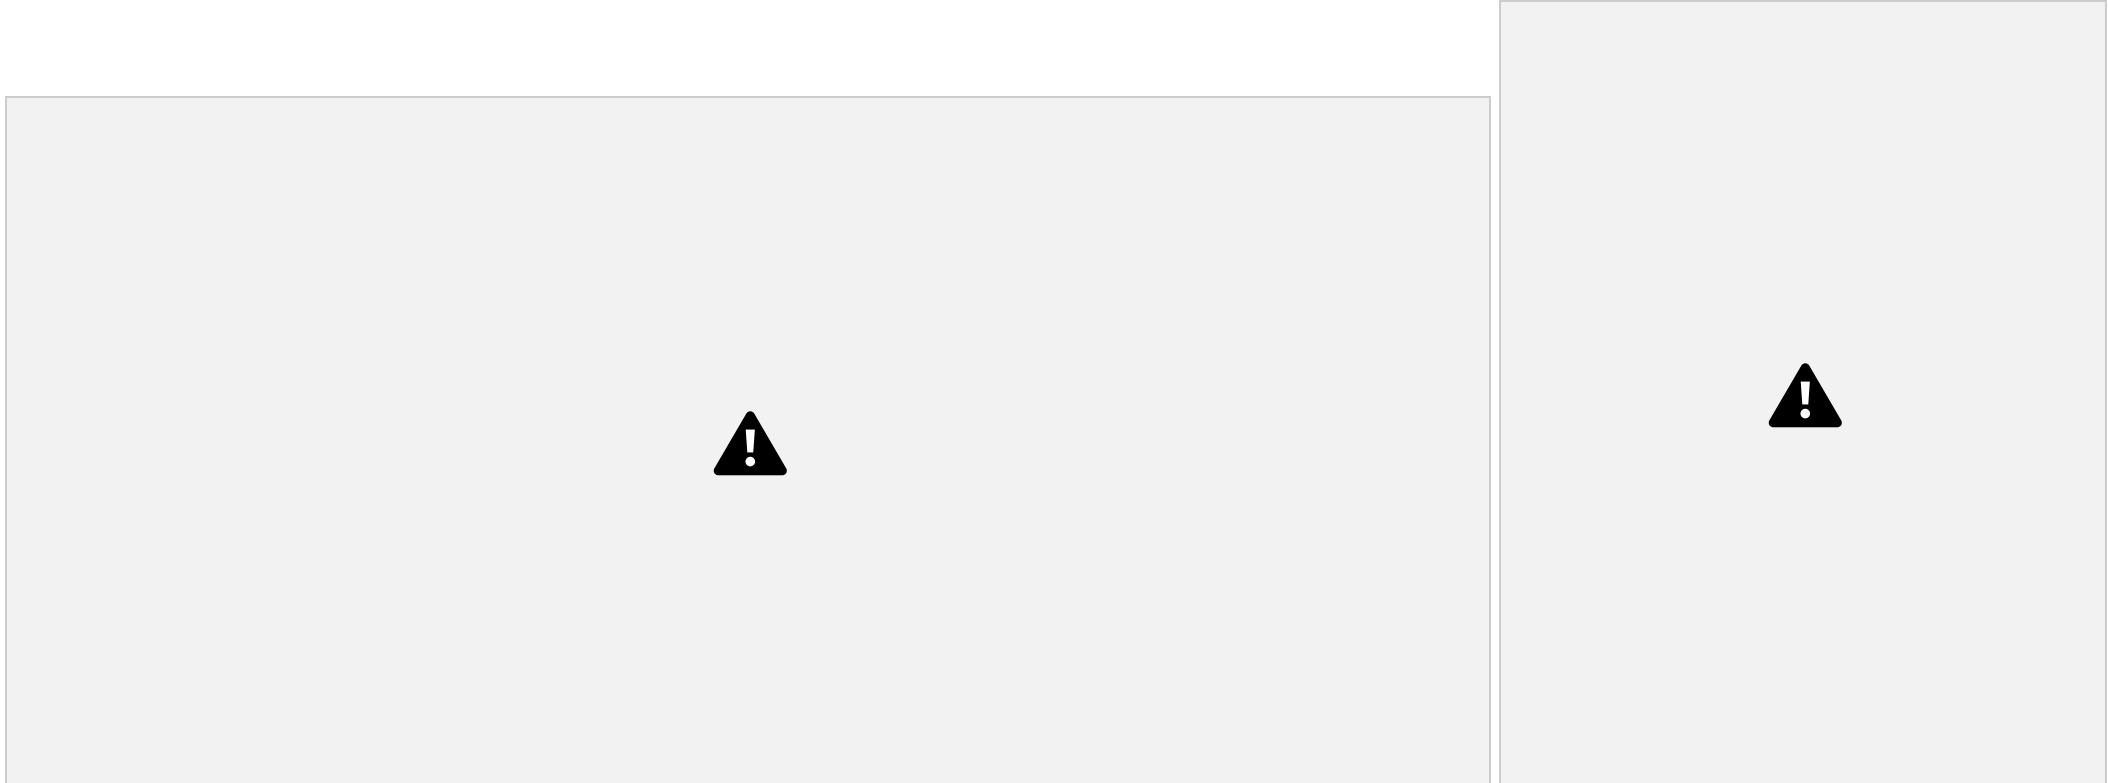
12

*McDuff, D., Schaekermann, M., Tu, T. et al. Towards accurate differential diagnosis with large language models. Nature (2025). <https://doi.org/10.1038/s41586-025-08869-4>*

## The “art” of prompting

- Prompting = a guiding question / instruction given to the model to shape the generated response

- Brittle!



Vinay, Rasita, et al. "Emotional Prompting Amplifies Disinformation

Generation in AI Large Language Models." *Frontiers in Artificial Intelligence* 8: 1543603.

- # What can you use an LLM for?
- Summarisation
  - Programming (Co-Pilot)
  - Question Answering
  - Reports from images
  - ...

- Recognise special pieces of information in a text ◦ “Named Entity Recognition”
- Sentiment analysis
- Fraud detection
- ...

Classification tasks

14

Generative tasks

# Key LLM training steps

2) Fine-tune

4) Validate / Evaluate 





# Tokenization

- Breaking down sentences into *tokens*
- The totality of tokens = the *vocabulary* of an LLM
  - Determined by the specific choice of *tokenizer*

- Tokens do *not* necessarily always correspond to *words*
  - *Why?*



## What is a Large Language Model? (Generalization)

- A (deep learning) model trained to **predict** the next **word** in a **sentence**
  - Token = words / amino-acids / genes / ....
  - What is a Language?
  - A **vocabulary** +

- **Sequences of tokens** that represent **information** in that language





## Pre-training: 2) embedding

- LLMs, like any machine learning system, work with *numbers*
- *Embeddings = a way to transform words into numbers, while preserving their semantics*



Source: <https://dkharazi.github.io/notes/ml/nlp/embedding><sub>21</sub>

Pre-training: Attention!











Source: [A](#)

[Multiscale Visualization of Attention in the Transformer Model](#), ACL 2019 Systems Demonstrations<sup>23</sup>

“Self-Supervised” / “Unsupervised” – but training data *matters*

- The illusion of self-supervision – actually a lot of effort for curating good data
  - Relies on work of thousands of human “data annotators”
  - Ethical issues: breaches of copyrighted material
  - Removing bias from training dataset
  - etc

Result: Foundation Model



Source: <https://viso.ai/deep-learning/foundation-models/>

## Non-exhaustive list of pitfalls and limitations...

- Hallucinations

- When it's wrong, it's confidently wrong!
- Cutoff date
  - An LLM has no knowledge past its training date

## **ANY SOLUTIONS?**

- Example: GPT 4.1 (April '25) has knowledge up to June '24
- COST!
  - 100 Million \$\$\$ to develop!
  - Environmental considerations...
- PRIVACY!
  - Please don't put sensitive information in ChatGPT...
- Lack of interpretability
- Lack of provenance / attribution
- Bias
- ....



## Open-source LLMs: Llama, Mistral...

Usually come in various size, the smaller the cheaper to run, but less “smart” ◆

■ ~**7B**: inference can run on your laptop in acceptable time if you have a GPU

◆ ■ ~**70B**: requires a larger server for decent inference time (1

A100/H100) ◆ ■ ~**400B**: largest, requires custom supercomputers (~4 H100s or more)

- Mixture-Of-Experts approaches such as Mixtral 8x7B or 8x22B combine multiple smaller models for better results with same hardware requirements
- HuggingFace: the most comprehensive resource to find open-source models
  - e.g. <https://huggingface.co/meta-llama/Meta-Llama-3-8B>



- Initially based on GPT3.5
- A “chatty” version of LLM
  - Based on Instruction DataSet (small, e.g. 50K etc) – Alpaca Style Prompt Template
- “Secret sauce”: Reinforcement Learning through Human Feedback (**RLHF**)

*We trained an initial model using supervised fine-tuning: **human AI trainers provided conversations in which they played both sides—the user and an AI assistant.** We gave the trainers access to model-written suggestions to help them compose their responses. We mixed this new dialogue dataset with the InstructGPT dataset, which we transformed into a dialogue format.*

Source: <https://openai.com><sup>29</sup>

## Reinforcement Learning through Human Feedback









30

RLHF – an ongoing journey...





Crowdsourcing expertise?

# Retrieval Augmented Generation (RAG)

- Complement the base model with *external knowledge base*

- This can be:

- PDFs
- Structured Data
  - CSVs
  - Relational DBs
  - ...
- Images
- ....



Source: <https://blog.gopenai.com>



# Fine-tuning a foundation model

- Requires a dataset of examples in the style of “instructions”, input and output



- Note: usually **expensive** to do





## Part II. From Theory (Closer To) Practice



How can you use LLMs to explore biodata?



- Directly ask questions (“zero-shot”)
- Formulate code (e.g. API calls) to answer questions
- Answering questions with context
  - Structured data (CSV files)
  - Unstructured data (PDF files)

 Ask LLM questions over biodata directly






“Cheating!”



Disable and try again!



Asking chatGPT question directly ("zero-shot")







Warning: in general, asking LLMs specialized questions directly will likely lead to hallucinated answers!







# Recommended: Ask Model To Ask an API







Recommended: Ask Model To Ask an API




Analysing local files?  
(LLMs with context)



Example: the SIB resources CSV


Download file from [here](#)






 Use ChatGPT for QA over CSV 

 Use ChatGPT for QA over CSV



Use ChatGPT for QA over CSV



Always



# Verify!

Using LLM for QA over PDF?





## Using ChatGPT for QA over complex PDF


Use an example ESMO clinical guideline, ask for a description of some of the figures or an analysis of a decision tree in one of the figures

E.g. [Early and locally advanced non-small-cell lung cancer \(NSCLC\): ESMO Clinical Practice Guidelines for diagnosis, treatment and follow-up](#)





# Question Answering over Structured Data






Expasy



# Knowledge Graphs curated by the SIB





# What is a Knowledge Graph?


- Graph data model
  - Easy to **interlink** and extend
  - Data usually stored as triples
  - Triple = <subject, predicate, object>
  - E.g. "*<gene X> isExpressedIn <anatomic entity Y>*"
- Can be queried using SPARQL:
  - Select \* where {  
    ?gene isExpressedIn ?anatEntity.  
}
- Play an important role in **Semantic Interoperability**



## Why are these useful?

- Explicit semantics through: -  
Standard identifiers
- Shared vocabularies
- Ontologies

Data Integration  
and FAIRification



- Ask complex questions that span across disciplines and across datasets

*What are the human genes involved in lung cancer  
with an ortholog expressed in the mouse?*

Research question: LLMs + KGs = ?



## Retrieval Augmented Generation (RAG)

Pan, Shirui, et al. "Unifying Large Language Models and Knowledge Graphs: A Roadmap." *arXiv preprint arXiv:2306.08302* (2023).

# How do we extract information from RDF data?



```
SELECT ?gene ?orthologous_protein2 WHERE {  
  SELECT * {  
    SERVICE <http://sparql.uniprot.org/sparql> {  
      SELECT ?protein1 WHERE {  
        ?protein1 a up:Protein;  
          up:organism/up:scientificName 'Homo sapiens' ;  
          up:annotation ?annotation .  
          ?annotation rdfs:comment ?annotation_text.  
          ?annotation a up:Disease_Annotation .  
          FILTER CONTAINS (?annotation_text, "lung cancer")  
        }  
      }  
    }  
  }  
  
  SERVICE <https://sparql.omabrowser.org/sparql/> {  
    SELECT ?orthologous_protein2 ?protein1 ?gene WHERE {  
      ?protein_OMA a orth:Protein .  
      ?orthologous_protein2 a orth:Protein .  
      ?cluster a orth:OrthologsCluster .  
      ?cluster orth:hasHomologousMember ?node1 .  
      ?cluster orth:hasHomologousMember ?node2 .  
      ?node2 orth:hasHomologousMember* [.....]  
      FILTER(?node1 != ?node2)  
    }  
  }  
  
  SERVICE <https://bgee.org/sparql/> {  
    ?gene genex:isExpressedIn ?anatEntity .  
    ?anatEntity rdfs:label 'lung' .  
    ?gene orth:organism ?org .  
    ?org obo:RO_0002162 taxon:10090 .}  
}
```

60

*What are the human genes involved in lung cancer with an ortholog expressed in the mouse?*



# LLMs for Biodata Exploration in Practice: ExpasyGPT



Writing SPARQL queries is hard and time-consuming. LLMs are great at it, but they need context

For complex questions finding the right context is not trivial.

We need a minimal structured way (metadata) to describe SPARQL endpoints.



# System architecture
















# Some Intermediate Conclusions...

- The LLM space is constantly evolving...
  - Tasks that seemed hard 6 months ago are now solved
  - Powerful interfaces to both structured and unstructured data...
  - ...but need to be adapted to your use case
- Careful when using (especially) closed-source LLMs
  - Don't use them for sensitive data!
- Infrastructure is a problem
  - Self-hosted (requires GPUs) or hosting platforms (paid)
  - Small models are (clearly) less performant
  - But big models are expensive to host

# References

- EMBL-EBI webinars on LLMs: <https://www.ebi.ac.uk/training/events/large-language-models-and-their-applications-bioinformatics/>
- ~~Developing an LLM: Building, Training, Fine-tuning~~ (relatively technical video): [https://www.youtube.com/watch?v=kPGTx4wcm\\_w](https://www.youtube.com/watch?v=kPGTx4wcm_w) - [Sebastian Raschka](#), author of “Build a Large Language Model (from scratch)”
- Detailed notebook for running LLMs for various text analyses (includes analysis of next word probability, embeddings generated etc): [https://colab.research.google.com/github/michael-franke/nlpNLG/blob/main/neural\\_pragmatic\\_nlg/07-LLMs/07b-pretrained-LLMs.ipynb](https://colab.research.google.com/github/michael-franke/nlpNLG/blob/main/neural_pragmatic_nlg/07-LLMs/07b-pretrained-LLMs.ipynb)
- Interactive overview of existing Language Models (from BERT to GPT3): [https://informationisbeautiful.net/visualizations/the-rise-of-generative-ai-large-language-models-llms-like-chatgpt/?trk=article-ssr-frontend-pulse\\_little-text-block](https://informationisbeautiful.net/visualizations/the-rise-of-generative-ai-large-language-models-llms-like-chatgpt/?trk=article-ssr-frontend-pulse_little-text-block)
- ~~Will we run out of data? Limits of LLM scaling based on human-generated data:~~ <https://arxiv.org/pdf/2211.04325>, interesting preprint analysing the limits of scaling beyond current model sizes
- Other LLMs to interact with: [Gemini](#), [Claude](#), ...
- ~~Language models for biological research: a primer~~, Nature Methods, '24
- ~~Extended LLM course with multiple hands-on notebook examples:~~ <https://github.com/mlabonne/llm-course>
- <https://github.com/ncbi-nlp/LLM-Medicine-Primer>
- ~~McDuff, D., Schaekermann, M., Tu, T. et al. Towards accurate differential diagnosis with large language~~

models. *Nature* (2025). <https://doi.org/10.1038/s41586-025-08869-4>

~~Prompting Amplifies Disinformation Generation in AI Large Language~~

models. *Frontiers in Artificial Intelligence* 8: 1543603.

- Vinay, Rasita, et al. "Emotional



# Thank you!

DATA SCIENTISTS FOR LIFE

**Ana-claudia.sima@sib.swiss**



## Part III. Hands-on: Setup

Install: uv, qdrant, get (Mistral) API key

